

# Agent Identity in ATB - Complete Guide

## How AI Agents Get SPIFFE Identities via SPIRE

ATB Team

January 2026

## Contents

<b>Agent Identity in ATB: Complete Guide</b>	<b>2</b>
Executive Summary . . . . .	2
<b>Part 1: How the Agent Gets Its First SPIFFE ID</b>	<b>2</b>
Overview . . . . .	2
Step 0: Admin Pre-Registers Workload Entry . . . . .	2
Step 1: Agent Workload Starts . . . . .	3
Step 2: Agent Requests SVID . . . . .	3
Step 3: SPIRE Agent Attests . . . . .	4
Step 4: X.509-SVID Issued . . . . .	4
<b>Part 2: Security Analysis</b>	<b>4</b>
Traditional vs SPIFFE Approach . . . . .	4
Attack Scenarios . . . . .	4
<b>Part 3: Different Attestation Methods</b>	<b>5</b>
Kubernetes (Most Common) . . . . .	5
Docker . . . . .	5
AWS . . . . .	5
Unix . . . . .	5
<b>Part 4: Complete Example</b>	<b>5</b>
1. Setup SPIRE . . . . .	5
2. Register Entry . . . . .	5
3. Create ServiceAccount . . . . .	5
4. Deploy Agent . . . . .	5
5. Agent Code . . . . .	6
<b>Summary</b>	<b>7</b>
<b>Related Documentation</b>	<b>7</b>

# Agent Identity in ATB: Complete Guide

## Executive Summary

This guide explains how AI agents in ATB (Agent Trust Broker) obtain their cryptographic identities using SPIFFE/SPIRE. Unlike traditional systems that use static API keys or secrets, ATB agents get short-lived, automatically-renewed certificates based on **where and how they run**.

## Key Benefits:

- No secrets to manage or rotate
  - Certificates valid for only 10 minutes
  - Identity tied to workload attestation
  - Private keys never touch disk
- 

## Part 1: How the Agent Gets Its First SPIFFE ID

### Overview

The identity issuance process has 4 main steps:

Step	Actor	Action
0	Admin	Registers workload entry in SPIRE Server
1	Kubernetes	Starts agent pod with namespace/SA
2	Agent	Calls SPIRE Workload API via Unix socket
3	SPIRE Agent	Attests workload and issues X.509 certificate

---

### Step 0: Admin Pre-Registers Workload Entry

Before any agent can get an identity, an admin registers an **entry**:

```
spire-server entry create \
  -spiffeID spiffe://example.org/agent/sales-bot \
  -parentID spiffe://example.org/spire/agent/k8s_psat/demo-cluster \
  -selector k8s:ns:ai-agents \
  -selector k8s:sa:sales-bot-sa
```

#### Parameters:

Parameter	Value	Meaning
-spiffeID	spiffe://example.org/agent/sales-bot	Identity to issue
-parentID	spiffe://... demo-cluster	Which SPIRE Agent can issue
-selector k8s:ns:ai-agents	Kubernetes namespace	Pod must be here
-selector k8s:sa:sales-bot-sa	ServiceAccount	Pod must use this

**This creates a rule:** “Any pod in namespace `ai-agents` using ServiceAccount `sales-bot-sa` gets the SPIFFE ID `spiffe://example.org/agent/sales-bot`”

---

## Step 1: Agent Workload Starts

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: sales-bot
  namespace: ai-agents
spec:
  template:
    spec:
      serviceAccountName: sales-bot-sa
      containers:
        - name: agent
          image: company/sales-bot:v1.2
          volumeMounts:
            - name: spire-agent-socket
              mountPath: /run/spire/sockets
              readOnly: true
      volumes:
        - name: spire-agent-socket
          hostPath:
            path: /run/spire/sockets
            type: Directory
```

---

## Step 2: Agent Requests SVID

Python:

```
from pyspiffe.workloadapi import X509Source

source = X509Source(
    workload_api_addr="unix:///run/spire/sockets/agent.sock"
)
svid = source.svid
print(f"My SPIFFE ID: {svid.spiffe_id}")
# Output: spiffe://example.org/agent/sales-bot
```

Go:

```
source, _ := workloadapi.NewX509Source(ctx,
    workloadapi.WithClientOptions(
        workloadapi.WithAddr("unix:///run/spire/sockets/agent.sock"),
    ),
)
```

```

svid, _ := source.GetX509SVID()
fmt.Printf("My SPIFFE ID: %s\n", svid.ID)

```

---

### Step 3: SPIRE Agent Attests

SPIRE Agent checks:

1. **WHO** is calling? (Unix socket peer credentials)
  2. **WHAT** Kubernetes context? (Query K8s API)
    - Namespace: ai-agents (MATCH)
    - ServiceAccount: sales-bot-sa (MATCH)
  3. **MATCH** against entries? YES
  4. **DECISION**: Issue SVID
- 

### Step 4: X.509-SVID Issued

Field	Value
Subject Alternative Name	URI: spiffe://example.org/agent/sales-bot
Validity	10 minutes (auto-renewed)
Public Key	Ed25519
Private Key	In memory only (never on disk)

---

## Part 2: Security Analysis

### Traditional vs SPIFFE Approach

Traditional	SPIFFE/SPIRE
API keys in env vars	No secrets to manage
Manual rotation	Auto-rotated every 10 min
Secrets can leak	Identity from attestation
Anyone with secret = access	Must be right workload

### Attack Scenarios

Attack	Why It Fails
Steal certificate	Expires in 10 minutes
Copy to another machine	Private key in memory only
Impersonate from different pod	Attestation checks namespace/SA
Create fake pod	Need K8s RBAC for namespace

---

## Part 3: Different Attestation Methods

### Kubernetes (Most Common)

```
-selector k8s:ns:ai-agents  
-selector k8s:sa:sales-bot-sa  
-selector k8s:pod-label:app:sales-bot
```

### Docker

```
-selector docker:label:ai.agent:sales-bot
```

### AWS

```
-selector aws:iamrole:arn:aws:iam::123456789:role/SalesBotRole
```

### Unix

```
-selector unix:uid:1001  
-selector unix:path:/opt/agents/sales-bot
```

---

## Part 4: Complete Example

### 1. Setup SPIRE

```
helm install spire-server spire/spire-server -n spire-system  
helm install spire-agent spire/spire-agent -n spire-system
```

### 2. Register Entry

```
kubectl exec -n spire-system spire-server-0 -- \  
spire-server entry create \  
-spiffeID spiffe://example.org/agent/sales-bot \  
-parentID spiffe://example.org/spire/agent/k8s_psat/demo-cluster \  
-selector k8s:ns:ai-agents \  
-selector k8s:sa:sales-bot-sa
```

### 3. Create ServiceAccount

```
apiVersion: v1  
kind: ServiceAccount  
metadata:  
  name: sales-bot-sa  
  namespace: ai-agents
```

### 4. Deploy Agent

```
apiVersion: apps/v1  
kind: Deployment
```

```

metadata:
  name: sales-bot
  namespace: ai-agents
spec:
  replicas: 1
  selector:
    matchLabels:
      app: sales-bot
  template:
    metadata:
      labels:
        app: sales-bot
    spec:
      serviceAccountName: sales-bot-sa
      containers:
        - name: agent
          image: company/sales-bot:v1.2
          env:
            - name: SPIFFE_ENDPOINT_SOCKET
              value: "unix:///run/spire/sockets/agent.sock"
      volumeMounts:
        - name: spire-agent-socket
          mountPath: /run/spire/sockets
          readOnly: true
      volumes:
        - name: spire-agent-socket
          csi:
            driver: "csi.spiffe.io"
            readOnly: true

```

## 5. Agent Code

```

from pyspiffe.workloadapi import X509Source
import os

def main():
    socket = os.environ.get(
        "SPIFFE_ENDPOINT_SOCKET",
        "unix:///run/spire/sockets/agent.sock"
    )
    source = X509Source(workload_api_addr=socket)

    svid = source.svid
    print(f"I am: {svid.spiffe_id}")
    # Output: spiffe://example.org/agent/sales-bot

```

---

## Summary

The agent **never stores secrets**. Its identity comes from:

- **Where it runs** (Kubernetes namespace)
- **How it runs** (ServiceAccount)
- **Cryptographic proof** (SPIRE attestation)

The entire process is automatic once the admin registers the workload entry.

---

## Related Documentation

- How AgentAuth Works
- SPIFFE Integration Guide
- Security Hardening
- Audit Logging