

How AgentAuth Works: Identity Validation and PoA Token Issuance

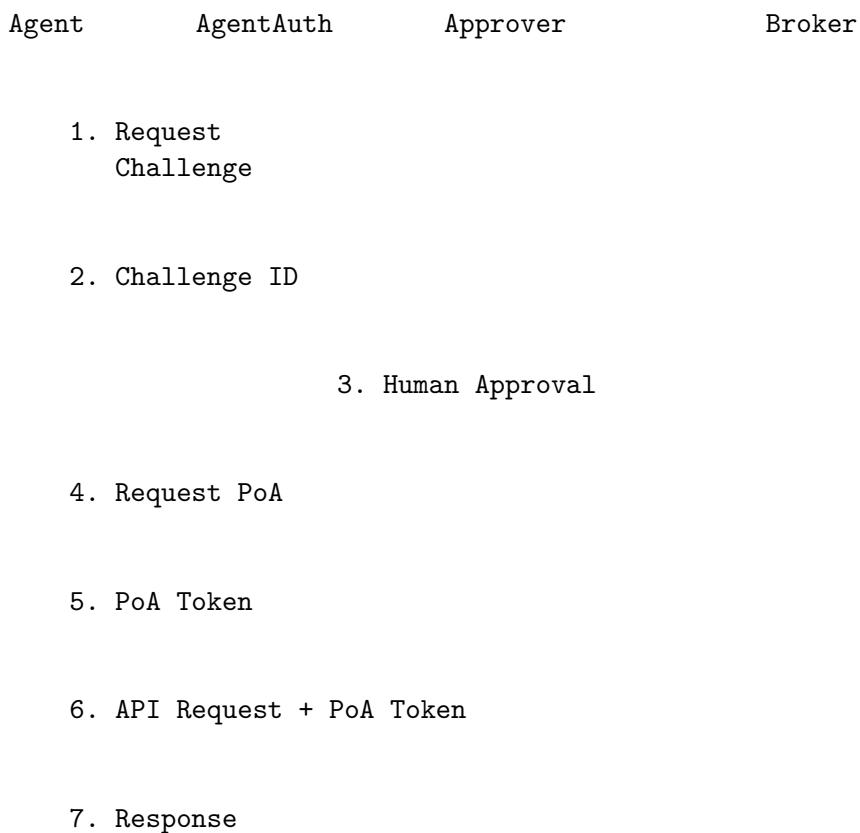
This document provides a step-by-step walkthrough of how ATB's AgentAuth service validates agent identity and issues Proof-of-Authorization (PoA) tokens.

Overview

AgentAuth is the identity and authorization gateway for AI agents. Before any agent can perform an action through ATB, it must:

1. Prove its identity (via SPIFFE/SPIRE)
2. Request authorization for a specific action
3. Obtain human approval (if required)
4. Receive a short-lived PoA token
5. Present the token to the Broker for action execution

AgentAuth Authorization Flow



Step 1: Agent Identity via SPIFFE

Before requesting any authorization, an agent must have a cryptographic identity issued by SPIFFE/SPIRE.

What is SPIFFE?

SPIFFE (Secure Production Identity Framework for Everyone) provides:

- **Automatic identity issuance** - No secrets to manage
- **Short-lived certificates** - Default 10-minute validity
- **Workload attestation** - Identity based on where code runs, not secrets

SPIFFE ID Format

Every agent has a SPIFFE ID that uniquely identifies it:

`spiffe://<trust-domain>/<workload-path>`

Examples:

- `spiffe://prod.company.com/agents/crm-assistant`
- `spiffe://prod.company.com/ns/ai-workloads/sa/clause-agent`

How the Agent Gets Its Identity

AI Agent
Workload

SPIRE Agent
(per node)

1. Request SVID via Workload API

SPIRE attests workload:

- Kubernetes namespace/SA
- Process UID/GID
- Docker labels

2. X.509-SVID (short-lived cert)

Certificate contains:

- SPIFFE ID as SAN URI
- 10-min validity
- Auto-rotated by SPIRE

AgentAuth Validates SPIFFE ID

When an agent connects to AgentAuth, its SPIFFE ID is validated:

```

// From agentauth/main.go - SPIFFE ID validation
var validSPIFFEIDRegex = regexp.MustCompile(
    `^spiffe://[a-zA-Z0-9]([a-zA-Z0-9.-]*[a-zA-Z0-9])?(/[a-zA-Z0-9._-]+)+$`)
)

func validateSPIFFEID(id string) error {
    if len(id) == 0 {
        return errors.New("SPIFFE ID is empty")
    }
    if len(id) > 2048 {
        return errors.New("SPIFFE ID too long (max 2048)")
    }
    if strings.Contains(id, "..") {
        return errors.New("SPIFFE ID contains path traversal")
    }
    if !validSPIFFEIDRegex.MatchString(id) {
        return errors.New("SPIFFE ID format invalid")
    }
    return nil
}

```

Step 2: Request a Challenge

The agent requests a challenge for a specific action. This creates an authorization request that must be approved before a PoA token is issued.

Challenge Request

```

POST /v1/challenge
Content-Type: application/json

{
    "agent_spiffe_id": "spiffe://prod.company.com/agents/crm-assistant",
    "act": "crm.contact.update",
    "con": {
        "max_records": 10,
        "allowed_fields": ["email", "phone"]
    },
    "leg": {
        "basis": "contract",
        "ref": "MSA-2026-001",
        "jurisdiction": "US",
        "accountable_party": {
            "type": "human",
            "id": "user@company.com"
        }
    }
}

```

}

What AgentAuth Validates

Field	Validation
agent_spiffe_id	Must match SPIFFE URI format, max 512 chars, no path traversal
act	Action name, max 256 chars, no null bytes
con	Constraints object, max 10 levels deep, no null bytes in keys/values
leg	Legal basis object, must include accountable_party

Rate Limiting

AgentAuth enforces rate limits to prevent abuse:

```
// Per-IP rate limiting (default: 100/min)
if !ipRateLimiter.Allow(clientIP) {
    return 429 Too Many Requests
}

// Per-agent rate limiting (default: 20/min)
if !agentRateLimiter.Allow(req.AgentSPIFFEID) {
    return 429 Too Many Requests
}
```

Challenge Response

```
{
    "challenge_id": "chal_x7k9m2...",
    "expires_at": "2026-01-15T10:05:00Z",
    "requires_dual_control": false,
    "approvers_needed": 1,
    "approval_hint": "POST /v1/approve with challenge_id and approver identity"
}
```

Step 3: Risk Assessment and Dual Control

AgentAuth determines if the action requires enhanced approval based on risk tier.

High-Risk Actions (Dual Control Required)

Certain actions automatically require two independent approvers:

```
// Default high-risk actions
highRiskActions := []string{
    "sap.vendor.change",
```

```

    "iam.privilege.escalate",
    "payments.transfer.execute",
    "ot.system.manual_override",
}

```

This can be configured via environment variable:

```
DUAL_CONTROL_ACTIONS="sap.vendor.change,payments.transfer.execute,..."
```

Explicit Dual Control

Requests can also explicitly require dual control:

```
{
  "leg": {
    "dual_control": {
      "required": true
    }
  }
}
```

Audit Log Entry

Every challenge creation is logged:

```
{
  "timestamp": "2026-01-15T10:00:00Z",
  "event": "challenge.created",
  "challenge_id": "chal_x7k9m2...",
  "agent_spiffe_id": "spiffe://prod.company.com/agents/crm-assistant",
  "action": "crm.contact.update",
  "risk_tier": "low",
  "requires_dual_control": false,
  "source_ip": "10.0.1.50",
  "success": true,
  "expires_at": "2026-01-15T10:05:00Z"
}
```

Step 4: Human Approval

A human approver must approve the challenge before a PoA token can be issued.

Approval Request

```
POST /v1/approve
Authorization: Bearer <approver-jwt>
Content-Type: application/json
```

```
{
```

```

        "challenge_id": "chal_x7k9m2..."
    }

```

Approver Authentication

AgentAuth supports multiple authentication methods:

Method	Configuration	Use Case
EdDSA JWT	APPROVER_ED25519_PUBLIC_KEY_PEM	Production (recommended)
RSA JWT	APPROVER_RSA_PUBLIC_KEY_PEM	Enterprise SSO integration
HMAC JWT	APPROVER_JWT_SECRET	Development/testing
Shared Secret	APPROVAL_SHARED_SECRET	Legacy (not recommended)

```

// JWT verification supports multiple algorithms
switch token.Method.Alg() {
    case "EdDSA":
        // Ed25519 signature verification
    case "RS256", "RS384", "RS512":
        // RSA signature verification
    case "HS256", "HS384", "HS512":
        // HMAC verification
}

```

Self-Approval Prevention

By default, the accountable party cannot approve their own requests:

```

if preventSelfApproval {
    accountableParty := req.getAccountablePartyID()
    if normalizeApproverID(approverID) == normalizeApproverID(accountableParty) {
        return "self-approval not allowed"
    }
}

```

Dual Control Approval

For high-risk actions requiring dual control:

1. First approver approves → Challenge status: 1/2 approvers
2. Second approver approves → Challenge status: 2/2 approvers (fully approved)
3. Both approvers must be distinct

```
{
    "challenge_id": "chal_x7k9m2...",
    "requires_dual_control": true,
    "approvers_needed": 2,
    "approvers_count": 1,
    "approvers": [
        {
            "id": "manager@company.com",

```

```

        "approved_at": "2026-01-15T10:01:00Z"
    }
],
"fully_approved": false
}

```

Step 5: PoA Token Issuance

Once the challenge is fully approved, the agent can exchange it for a PoA token.

Token Request

POST /v1/token
Content-Type: application/json

```
{
  "challenge_id": "chal_x7k9m2..."
}
```

Token Generation

AgentAuth generates a signed JWT with specific claims:

```
claims := PoAClaims{
    Act: challenge.Req.Act,           // Authorized action
    Con: challenge.Req.Con,          // Constraints
    Leg: challenge.Req.Leg,          // Legal basis
    RegisteredClaims: jwt.RegisteredClaims{
        Issuer:      "atb-agentauth",
        Subject:     challenge.Req.AgentSPIFFEID,
        Audience:   []string{"atb-broker"},
        ExpiresAt:   jwt.NewNumericDate(now.Add(poaTTL)),
        IssuedAt:   jwt.NewNumericDate(now),
        ID:          mustRandID("poa_"),
    },
}
```

Token Signing (EdDSA)

Tokens are signed with Ed25519 for security and performance:

```
token := jwt.NewWithClaims(jwt.SigningMethodEdDSA, claims)
token.Header["kid"] = kid // Key ID for rotation support

signedToken, err := token.SignedString(privateKey)
```

PoA Token Structure

Header:

```
{
  "alg": "EdDSA",
  "typ": "JWT",
  "kid": "abc123..." // Key ID for verification
}

Payload:
{
  "iss": "atb-agentauth",
  "sub": "spiffe://prod.company.com/agents/crm-assistant",
  "aud": ["atb-broker"],
  "exp": 1736936700, // 5 minutes from issuance
  "iat": 1736936400,
  "jti": "poa_xyz789...",
  "act": "crm.contact.update",
  "con": {
    "max_records": 10,
    "allowed_fields": ["email", "phone"]
  },
  "leg": {
    "basis": "contract",
    "ref": "MSA-2026-001",
    "accountable_party": {
      "type": "human",
      "id": "user@company.com"
    }
  }
}
```

Token Response

```
{
  "poa_token": "eyJhbGciOiJFZERTQSIs...",
  "expires_at": "2026-01-15T10:05:00Z",
  "token_id": "poa_xyz789..."
}
```

Step 6: Key Rotation Support

AgentAuth supports seamless key rotation for production environments.

Multiple Signing Keys

```
# Primary key (used for signing)
POA_SIGNING_ED25519_PRIVKEY_PEM="..."

# Previous key (for verification during rotation)
```

```

POA_SIGNING_ED25519_PRIVKEY_PEM_PREV="..."  

# Next key (pre-staged for upcoming rotation)  

POA_SIGNING_ED25519_PRIVKEY_PEM_NEXT="..."

```

JWKS Endpoint

The Broker verifies tokens using the JWKS endpoint:

```

GET /.well-known/jwks.json
{
  "keys": [
    {
      "kty": "OKP",
      "crv": "Ed25519",
      "use": "sig",
      "alg": "EdDSA",
      "kid": "abc123...",
      "x": "base64url-encoded-public-key"
    },
    {
      "kty": "OKP",
      "crv": "Ed25519",
      "use": "sig",
      "alg": "EdDSA",
      "kid": "def456...",
      "x": "base64url-encoded-previous-key"
    }
  ]
}

```

Step 7: Broker Token Verification

When the agent presents the PoA token to the Broker, it is validated before proxying the request.

Broker Verification Steps

1. **Parse JWT** - Extract header, payload, signature
2. **Verify Signature** - Using JWKS from AgentAuth
3. **Check Expiration** - Token must not be expired
4. **Validate Audience** - Must include “atb-broker”
5. **Match SPIFFE ID** - Token `sub` must match client certificate
6. **Verify Action** - Requested action must match `act` claim
7. **Enforce Constraints** - Request must satisfy `con` claims
8. **OPA Policy** - Query OPA for additional policy checks

OPA Policy Evaluation

```
# From opa/policy/poa.rego
default allow := false

allow if {
    valid_signature
    not_expired
    valid_audience
    valid_action
    constraints_satisfied
}
```

Security Controls Summary

Control	Implementation
Identity	SPIFFE/SPIRE X.509-SVIDs
Input Validation	Regex, length limits, null byte checks
Rate Limiting	Per-IP and per-agent limits
Token TTL	Max 900 seconds (15 min), default 300s
Signing	Ed25519 (EdDSA) with key rotation
Self-Approval Prevention	Accountable party cannot approve
Dual Control	Two approvers for high-risk actions
Audit Logging	Structured JSON for all events
Security Headers	X-Content-Type-Options, CSP, etc.

Configuration Reference

AgentAuth Environment Variables

Variable	Default	Description
LISTEN_ADDR	:9090	HTTP listen address
POA_ISSUER	atb-agentauth	Token issuer claim
POA_TTL_SECONDS	300	Token validity (max 900)
CHALLENGE_TTL_SECONDS	300	Challenge validity (max 900)
RATE_LIMIT_PER_IP	100	Requests per minute per IP
RATE_LIMIT_PER_AGENT	20	Requests per minute per agent
REQUIRE_JWT_AUTH	false	Require JWT for approvers
ALLOW_SELF_APPROVAL	false	Allow self-approval (not recommended)
DUAL_CONTROL_ACTIONS	(defaults)	Comma-separated high-risk actions
POA_SIGNING_ED25519_PRIVKEY(gPEM)		Primary signing key

Troubleshooting

Common Issues

“SPIFFE ID format invalid” - Ensure SPIFFE ID matches: `spiffe://<domain>/<path>` - No path traversal (...), no special characters

“rate limit exceeded” - Check RATE_LIMIT_PER_IP and RATE_LIMIT_PER_AGENT settings - Implement exponential backoff in agent

“self-approval not allowed” - The accountable party in `leg.accountable_party.id` cannot approve - Use a different approver or set `ALLOW_SELF_APPROVAL=true` (not recommended)

“challenge expired” - Challenges expire after CHALLENGE_TTL_SECONDS (default 5 min) - Request a new challenge

“JWT verification failed” - Check approver JWT configuration - Verify issuer is in APPROVER_JWT_ISSUERS list - Ensure public key matches the signing key

Related Documentation

- Authentication & Authorization Guide
- SPIFFE/SPIRE Integration Guide
- Security Hardening Guide
- Audit Logging
- API Reference